

COMP5046: Classification with Features

Joel Nothman

`joel.nothman@sydney.edu.au`

School of Information Technologies
University of Sydney

2018-04-10

Part I

Naïve Bayes

What are we trying to do?

- some examples:
 - try to predict if it will rain tomorrow
 - try to predict what someone will say next
 - classify texts into categories, e.g. law, health, etc.
- we have a lot of information to exploit in language
- **build a model from training data that can predict future data**

Outline

- Bayes' Rule for classification
- Naïve Bayes
- **Feature representation**
- Perceptron models
- Maximum Entropy models

What are we trying to calculate?

- x is the information we have (the observed context):
 - who is talking and what they've already said
 - the entity and their Wikipedia article
- y is what we're trying to predict:
 - what word they'll say next
 - the category the entity fits into

• now calculate $p(y|x) \forall y$ (all classes)

• and select the y that maximises $p(y|x)$, formally:

$$\underset{y}{\operatorname{argmax}} p(y|x) \quad (1)$$

Bayes' Rule derives from conditional probability

- conditional probability leads to the chain rule:

$$p(y|x) \equiv \frac{p(y \cap x)}{p(x)} \quad (2)$$

$$p(y \cap x) = p(x)p(y|x) = p(y)p(x|y) \quad (3)$$

- rewriting Equation 3 we get **Bayes' rule** (or theorem):

$$p(y|x) \equiv \frac{p(y)p(x|y)}{p(x)} \quad (4)$$

Bayes' Rule splits $p(y|x)$ into pieces

- combine the last two slides together
- we can calculate the most likely y given x :

$p(x)$ is constant for every instant

$$\operatorname{argmax}_y p(y|x) = \operatorname{argmax}_y \frac{p(y)p(x|y)}{p(x)} \quad (5)$$

$$= \operatorname{argmax}_y p(y)p(x|y) \quad (6)$$

- we can drop out the denominator for fixed x when we only want y
- **often useful when $p(x|y)$ and $p(y)$ easier to calculate**
- **or they encode the structure of the problem**
- NB: $p(x|y)$ can cause confusion, since it seems back to front

Calculating the prior probability, $p(y)$, is usually easy

- how many times the class y occurs divided by the total number of training instances

$$p(y) = \frac{\text{count}(y)}{\sum_{y_i \in Y} \text{count}(y_i)} \quad (7)$$

- **relative frequency calculated over our training data**
- smoothing is still important for rare classes

Calculating $p(x|y)$ is often difficult

- conditioning on y is easy in classification
- but modelling the probability of $x \dots$
- $p(x|y)$ can also be calculated using relative frequency if you assume x can be treated as an atom (a unit)
- although smoothing is still critical
- **difficult to calculate if x is a complex, structured object**
e.g. a Wikipedia article
- probabilities for $p(x|\text{anything})$ become vanishingly small
- could use a language model for text, but we may want to encode diverse information

Features identify aspects of a context

- features describe a particular state of the decision context
- **and** the hypothesised/actual class for classification tasks

| | |
|--|---------------------|
| (actress in Wikipedia text, Film/TV) | Jessica Alba |
| (murderer in Wikipedia text, Crime) | Mark “Chopper” Reid |
| (song in Wikipedia text, Music) | Lady Gaga |

- features don't need to be absolute predictors of a class
- **i.e. features can be ambiguous**
- they can also be *negative evidence* for a class

A feature is...

- often a boolean *indicator* function (i.e. returning 0 or 1)

$$f_i(x, y) = \text{attribute}(x) \cap (\text{class} = y) \quad (8)$$

- where $\text{attribute}(x)$ checks evidence from the context
 - might identify a boolean property (e.g. `islower(word)`)
 - might pick out elements of the context (e.g. `prev_word == Mr`)
- `islower` and `prev_word` are *contextual predicates*
- `Mr` is the *value* of the `prev_word` contextual predicate
- more generally, f_i is a (bounded) real-valued function
- and it can pick out aspects of the class y , as well as the context x

A feature is often defined in terms of x only

- making feature vector a function of x and y allows y to be structured
 - there may be too many classes to learn separate parameters for each
 - similar y can share features and hence learnt parameters
- for simple classification, we ignore structure of y
 - use same attributes of x for features with each class y

| | PER | | | | LOC | | | | ORG | | | |
|------------------------------|-----|---|---|---|-----|---|---|---|-----|---|---|---|
| $f(\text{John}, \text{PER})$ | 1 | 0 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| $f(\text{John}, \text{LOC})$ | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 1 | 0 | 0 | 0 | 0 |
| $f(\text{John}, \text{ORG})$ | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 1 |

- note y here does not need to be the *gold standard* class for x
 - we can calculate feature values for each y being considered for some x

Features in Naïve Bayes

- split the context (training instance) into individual features:
- replace x with $x = \{x_1, x_2, \dots, x_n\}$

$$p(x|y) = p(x_1, \dots, x_n|y) \quad (9)$$

$$= p(x_1|y)p(x_2|x_1, y) \dots p(x_n|x_{n-1} \dots x_1, y) \quad (10)$$

- this is no easier to calculate than before!
- **assume conditional independence between features to simplify:**

$$p(x|y) = p(x_1|y)p(x_2|y) \dots p(x_n|y) \quad (11)$$

$$= \prod_i p(x_i|y) \quad (12)$$

Reminder: Conditional Independence

- two events A and B are **conditionally independent** iff:

$$p(A \cap B | C) = p(A | C)p(B | C) \quad (13)$$

- i.e. $p(A | B \cap C) = p(A | C)$ and $p(B | A \cap C) = p(B | C)$
- knowing C gives no knowledge about $A \cap B$ beyond what C indicates individually about A or B
- i.e. knowing C may give some knowledge about A or B

Complex Features

- features can be arbitrarily complex
 - e.g. document level features in named entity recognition
(document = cricket & current word = Lancashire, **ORG**)
⇒ hopefully tag Lancashire as **ORG** not **LOC**
- features can be combinations of atomic features
 - (current word = David & next word = Jones, **ORG**)
⇒ hopefully tag David as **ORG** not **PER**
- **feature engineering** during model development to define useful features
 - designed to target frequent errors

Conditional independence is a lie!

- conditional independence is an assumption of Naïve Bayes
- **unfortunately the assumption is usually totally wrong**
- for example:
unigram=actor in text, `Film/TV`
is not independent of
bigram=American actor in text, `Film/TV`
- **often we get away with it anyway!**

Estimation of Naïve Bayes parameters

- each $p(x_i|y)$ is calculated using relative frequencies (see *lab!*):

$$p_{iy} = \frac{\text{count}(\text{feature } i \text{ in class } y)}{\sum_{i'} \text{count}(\text{feature } i' \text{ in class } y)} \quad (14)$$

of feature i in class / (sum of any feature in class)

- parameters*: values calculated in training, used to predict
- this relative frequency is the Maximum Likelihood Estimate**
- MLE is one approach to *faithful representation*
- what about previously unseen features?**
- smoothing: redistribute mass to unseen features for more *sensible generalisation***

Part II

Perceptron models

Generative and Discriminative Models

- Using Bayes' Rule requires a model of an observation x
 - a **generative model** calculates $P(x, y)$
 - Can be hard to design a principled model
i.e. without naïve independence assumptions
- To predict the best y , we can learn $P(y|x)$ directly
 - or even just $\text{score}(y|x)$ (not a valid probability)
 - a **discriminative model**
 - State of the art for many tasks
 - The system can efficiently learn how to trust features

Looking at features differently

- instead of thinking of features as events in a joint probability:

$$p(x|y) = p(x_1, \dots, x_n|y) \quad (15)$$

- think of them as being weighted by how likely they are
- e.g. for named entity recognition:
 - (word = John, PER) \implies high weight
 - (word = John, LOC) \implies negative weight
 - (word has an even number of letters, PER) \implies near-zero weight

Feature vectors

classifying a single word

- for all **active** features, add up all the associated weights

| | PER | | | | LOC | | | | ORG | | | |
|---------------------------------------|-----|---|---|---|-----|---|----|---|-----|---|---|---|
| $\mathbf{f}(\text{John}, \text{PER})$ | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| $\mathbf{f}(\text{John}, \text{LOC})$ | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
| \mathbf{w} | 3 | 1 | 1 | 0 | -2 | 2 | -1 | 0 | -1 | 1 | 1 | 0 |

$$\text{score}(x, y, \mathbf{w}) = \mathbf{w} \cdot \mathbf{f}(x, y) \quad (16)$$

$$= \sum_i w_i f_i(x, y) \quad (17)$$

- dot product between the weight vector and feature vector

$$\text{score}(\text{John}, \text{PER}, \mathbf{w}) = (1 \times 3) + (1 \times 1) = 4 \quad (18)$$

$$\text{score}(\text{John}, \text{LOC}, \mathbf{w}) = (1 \times -2) + (1 \times -1) = -3 \quad (19)$$

Weights can be hard to interpret

- can incorporate many complex overlapping features
 - no independence assumptions
- allows linguistically motivated features
- but can be hard to interpret:
 - $x = \text{the cats sat on the mat}$
 - $f_{(\text{has word cats})}(x) = f_{(\text{has stem cat})}(x) = 1$
 - $w_{(\text{has word cat})} = 1$
 - $w_{(\text{has word cats})} = 1$
 - $w_{(\text{has stem cat})} = 1$

Imagine we want to know which features are most important. . .

Weights can be hard to interpret

- can incorporate many complex overlapping features
 - no independence assumptions
- allows linguistically motivated features
- but can be hard to interpret:

- $x = \text{the cats sat on the mat}$
- $f_{(\text{has word cats})}(x) = f_{(\text{has stem cat})}(x) = 1$
- $w_{(\text{has word cat})} = 1$
- $w_{(\text{has word cats})} = 1$
- $w_{(\text{has stem cat})} = 1$

Imagine we want to know which features are most important. . .
Weight is effectively double how it appears

Weights can be hard to interpret

- can incorporate many complex overlapping features
 - no independence assumptions
- allows linguistically motivated features
- but can be hard to interpret:

discriminative models are hard to interpret since the weights do not provide any information

- $x = \text{the cats sat on the mat}$
- $f_{(\text{has word cats})}(x) = f_{(\text{has stem cat})}(x) = 1$
- $w_{(\text{has word cat})} = 1$
- $w_{(\text{has word cats})} = 1$
- $w_{(\text{has stem cat})} = -1$

Weight is effectively much smaller than it appears

Optimisation

- predicted class for x is: $\underset{y}{\operatorname{argmax}} \mathbf{w} \cdot \mathbf{f}(x, y)$
 - prediction derives from a **linear** function of features
 - (other discriminative models may be non-linear)
- weights \mathbf{w} are parameters of the model
- need to find \mathbf{w} that gives good predictive performance
 - faithful representation of training data
 - sensible generalisation to new observations
- many approaches to finding \mathbf{w}
 - perceptron: simple; online
 - maximum entropy: principled; probabilistic interpretation
 - support vector machines
 - ...

Perceptron

- the goal is to minimise training errors
 - **online** – uses decoding/prediction as part of the training process
-
- perceptron is not often used in state of the art systems
 - but it gives you an intuition around optimising a linear model
 - and will be used for Assignment 3 where an online learner is needed

The algorithm

- set weights $\mathbf{w}^{(0)}$ to 0
- for t iterations:
 - for each training instance x , with correct class y :
 - classify x with the current weights $\mathbf{w}^{(t)}$ to find y^* (y^* is the best guess, y^* will need to be chosen randomly as $w=0$)
 - if $y = y^*$, do nothing
 - otherwise, update weights: $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \mathbf{f}(x, y) - \mathbf{f}(x, y^*)$ update weights, former weight + features of true class - features of current predicted class
- y is the correct, gold-standard class from the training data
- y^* is the current model's best guess at the class
- If these are not the same, we try to push the weights towards a better prediction given the same features:
 - inhibit the model choosing y^*
 - encourage the model to choose y

Updating weights

- Three cases for each feature:
 - it's active in correct class but not in the model's class
 - it's active in model's class but not in the correct class
 - it's the same in both classes

| | | | | |
|----------------------|---|----|-----|---|
| $\mathbf{w}^{(t)}$ | 0 | -1 | ... | 2 |
| $\mathbf{f}(x, y)$ | 1 | 0 | ... | 0 |
| $\mathbf{f}(x, y^*)$ | 0 | 1 | ... | 0 |
| $\mathbf{w}^{(t+1)}$ | 1 | -2 | ... | 2 |



Perceptron

- decoding:


$$y^* = \underset{y}{\operatorname{argmax}} \mathbf{w} \cdot \mathbf{f}(x, y) \quad (20)$$

- simple and fast to train
- problems with overtraining

Separability

learning a weight vector = learning a line between the classes

for non-separable data, the weights continually move around



<http://www.cs.berkeley.edu/~klein/papers/classification-tutorial-naacl2007.pdf>

Averaged perceptron

- set weights $\mathbf{w}^{(0)}$ to 0
- for t iterations:
 - for each training instance x , with correct class y :
 - classify x with the current weights $\mathbf{w}^{(t)}$ to find y^*
 - if $y = y^*$, do nothing
 - otherwise, update weights: $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \mathbf{f}(x, y) - \mathbf{f}(x, y^*)$
 - save $\mathbf{w}^{(t)}$
- **calculate the average weights over all $\mathbf{w}^{(t)}$**
- prediction on new data:

$$\operatorname{argmax}_y \sum_t \mathbf{w}^{(t)} \cdot \mathbf{f}(x, y) \quad (21)$$

Part III

Maximum Entropy Models

Maximum Entropy principle

(~ stats: Logistic Regression or
NN: soft max activation with soft entropy loss)

- *Maximum Entropy* modelling:
 - predicts observations from training data (faithful representation)
 - this **does not uniquely identify the model**
- chooses the model which has the most uniform distribution:
 - i.e. the **model with the maximum entropy** (sensible generalisation)

The maximum entropy formulation: constrained optimisation

- Find a model $p(y_i|x_i)$ which
- Is constrained for each feature j :
the expectation under the model equals the empirical expectation

$$E_p f_j = E_{\tilde{p}} f_j \quad \forall \text{ features } f_j \quad (22)$$


$$\dots \quad (23)$$

$$\sum_i p(y_i|x_i) f_j(x_i, y_i) = \sum_i f_j(x_i, y_i) \quad (24)$$

- Maximises entropy $H(p) = -\sum_x p(x) \log_2 p(x)$
 - a measure of uncertainty (\therefore uniformity) in a distribution

Maximising entropy under constraints

constraint = max entropy falls on the plane



[http://www.cs.berkeley.edu/~klein/papers/
maxent-tutorial-slides.pdf](http://www.cs.berkeley.edu/~klein/papers/maxent-tutorial-slides.pdf)

Logistic regression finds the same weights

- MaxEnt is hard to optimise directly
- But results in the same model as **logistic regression**

Deriving logistic regression

- we want the linear scoring function we saw before

$$y^*(x) = \operatorname{argmax}_y \sum_t \mathbf{w}^{(t)} \cdot \mathbf{f}(x, y) \quad (25)$$

- we want the model to be probabilistic
- how can we calculate the weights?
- maximise the likelihood of the training data (MLE)

$$L(\mathbf{w}) = \prod_i p(y_i | x_i, \mathbf{w}) \quad (26)$$

we want to find the weight vector \mathbf{w} which gives the max probability

Weights

- turn the score function we had previously into a probability distribution

$$p(y|x, \mathbf{w}) = \mathbf{w} \cdot \mathbf{f}(x, y) = \sum_i w_i f_i(x, y) \quad (27)$$

but there are negative weights and we can't have negative probabilities

$$p(y|x, \mathbf{w}) = \exp(\mathbf{w} \cdot \mathbf{f}(x, y)) \quad (28)$$

but a probability distribution must sum to 1

$$p(y|x, \mathbf{w}) = \frac{\exp(\mathbf{w} \cdot \mathbf{f}(x, y))}{\sum_{y'} \exp(\mathbf{w} \cdot \mathbf{f}(x, y'))} \quad (29)$$

Maximising the likelihood

- easier to calculate using the log-likelihood:

$$L(\mathbf{w}) = \log \prod_i p(y_i | x_i, \mathbf{w}) \quad (30)$$

$$= \log \prod_i \frac{\exp(\mathbf{w} \cdot \mathbf{f}(x_i, y_i))}{\sum_{y'} \exp(\mathbf{w} \cdot \mathbf{f}(x_i, y'))} \quad (31)$$

$$= \sum_i \log \left(\frac{\exp(\mathbf{w} \cdot \mathbf{f}(x_i, y_i))}{\sum_{y'} \exp(\mathbf{w} \cdot \mathbf{f}(x_i, y'))} \right) \quad (32)$$

$$= \sum_i \left(\mathbf{w} \cdot \mathbf{f}(x_i, y_i) - \log \sum_{y'} \exp(\mathbf{w} \cdot \mathbf{f}(x_i, y')) \right) \quad (33)$$

- this is an unconstrained optimisation problem**

Convexity

- the surface is convex, because the functions in the sum are convex

$$H(p) = - \sum_{x,y} p(y|x) \log_2 p(y|x) \quad (34)$$

$$L(\mathbf{w}) = \sum_i \left(\mathbf{w} \cdot \mathbf{f}(x_i, y_i) - \log \sum_{y'} \exp(\mathbf{w} \cdot \mathbf{f}(x_i, y')) \right) \quad (35)$$

- no local maxima**
- this makes the search problem simpler
- e.g. can use gradient descent

Regularization (or smoothing)

- many of the model's weights are optimised based on very little data
 - there are a large number of weights to set
 - the training data can be very sparse
- this leads to a model that is **overfitted** to the training data
- it will perform poorly on unseen data
- Solution: assume a **Gaussian prior** distribution over weights
 - equivalently: add L_2 norm of weight matrix to objective function
 - **L_2 regularisation**
 - penalises weights with large magnitude
 - downplays highly discriminative but unreliable (low freq.) features

Generative and Discriminative Models

- **joint** or **generative** models
 - model **both** observed instance (i.e. outcome) and classification
 - as if class **generated** instance
 - requires modelling probability of an observation
 - language models, Naïve Bayes, hidden markov models, PCFGs
 - trained using (smoothed) Maximum Likelihood Estimate
- **conditional** or **discriminative** models
 - model classification **given** observed instance
 - Maximum Entropy, SVMs, perceptrons
 - trained using (smoothed) conditional Maximum Likelihood Estimate

Other Discriminative Models

- Maximum Entropy models weight individual evidence (features)
- other approaches, e.g. SVMs or perceptrons weight evidence
- although all perform classification, other approaches are harder to interpret as probability distributions over classes
- important when uncertainty is important downstream

Things to do with a model (h/t UMass)

| Task | Given | Find |
|-------------------------|--|--------------------------------|
| Engineer features/model | $\langle x, y \rangle_i$ | $\mathbf{f}(x, y, \mathbf{w})$ |
| Estimate parameters | $\mathbf{f}(x, y, \mathbf{w}), \langle x, y \rangle_i$ | \mathbf{w} |
| Decode/predict | $\mathbf{f}(x, y, \mathbf{w}), x, \mathbf{w}$ | y^* for x |

Take away

- Feature representation of text
- Feature engineering
- Generative vs discriminative models
- Linear weighting of features
- Perceptron algorithm and averaged/voted variants
- Terms: parameters, estimation, decoding
- Again: faithfulness vs generalisation